

Notes on an Intro to Industrial Programming

As Illustrated by an LZW Implementation

Aaron Gorenstein

July 21, 2024

I'm exploring using more advanced projects for empowering post-baccelelaurette self-studies. This is one such example. I have (had) grand plans to make *this very document* a resource to that end. Practically speaking, however, I think I need to do this more first. So this is more a collection of neat thing I'd want to use this LZW implementation to talk about.

Contents

1. Introduction	2
2. Design Overview	2
3. Implementation Details	5
3.1. Reading and Writing Bits and Bytes	5
3.2. Interlude: Modules versus Objects	7
3.3. State-Machine Based Dictionary	7
3.4. Data Structure Design	7
3.5. Feature: Clear Codes	9
3.6. Using our Features: EMA	9
3.7. Debugging Support	9
4. Testing: Fuzzing	9
5. Performance	9
5.1. Compression Ratio	10
5.2. Throughput	10
6. What's Next?	10
7. Industrial Comparison	10

A. Core Library Implementation	10
B. Command Line Implementation	21

1. Introduction

What Does This Give You? This is *not* going to be an introduction to how LZW works: there are plenty of good resources for that. This is a resource along two dimensions:

1. Those who feel like they understand that introduction, but have questions about what implementation “really looks like” as you get to the nitty-gritty of an algorithm;
2. Those with an undergraduate or introductory background in programming and are curious to see how a homework-level implementation (as might be found on Rosetta code) can be taken to a more professional level.

Why Trust Me? I have no meaningful prior experience in understanding or implementing compression algorithms. I have a few years of graduate study in computational complexity theory (so I have some background on the formal, mathematical side of things), I spent the better part of a decade on the Microsoft C++ compiler (so I’m familiar with applying more formal computer science to industry, as well as industrial C-or-C++ programming), and most recently am trying out distributed systems at MongoDB (so I have my finger on the pulse of cool hip things). I have a great interest in teaching computer science, which I’ve pursued throughout my education and now with intermittent volunteer opportunities (some of which were quite extensive).

The end result of the code here will *not* be competitive with, say, **gzip**. However, it will employ a considerable number of performance-focused implementation and design decisions. In this way it can be understood as an intermediate step between homework and “real program”. Having not implemented **gzip** (which, to be clear, has a different algorithm), of course, to a degree that this is a useful stepping stone is taken on faith. However, the actual design decisions I’ll highlight (various **struct** design decisions, for instance) I’ve all used in various “real” production code.

Why Write This? I enjoy studying computer science, in particular exploring the boundary between theory, academia, and industry, and sharing what I learn. This write-up is partly my own notes so I feel like I will “forever understand” LZW to a deep degree, as well as an avenue to talk about some less common programming tricks in-situ.

2. Design Overview

Let’s review how LZW-encoding works:

- We’re receiving a stream of bytes, and wish to emit a stream of bytes.

- As we read in our input, we’re constantly finding the longest-matching-prefix that we’ve seen-so-far.
- Once we no longer have a previously-seen-prefix, we map this *new* string (which differs from a previously-seen-prefix by the single, additional, final byte) to a new integer key.
- We emit the key for the *previous* prefix.
- So, that is, we *don’t* emit the just-created new key: we instead restart the prefix-match process, starting with that final byte we used last time. (Restarting with that newly-seen character is how the decoder is able to know which character it was.)
- Available features (complications) include bounding the number of prefixes we save in our dictionary, resetting our dictionary (to help in cases where the input dramatically changes character “halfway through” compression), and variable-width encoding of keys.

An apparently-canonical implementation can be found on the internet.

I’ll introduce our implementation by comparison with the (single?) design that seems to pervade web resources. The listing in fig. 1, taken from the web-resources for Sedgwick¹ seems to express the common implementation. Of course it accurately implements (in this case) the simpler LZW variant with fixed-width (in this case, 12 bits) encoding. However, resources (mainly Wikipedia, but also that same Sedgwick site and the linked-to chapter on compression, here, and maybe even the original paper?) talks about variable-width.

To be clear, this is obviously a fine implementation to learn LZW—and the more general algorithmic considerations of compression—with. My presentation here is better understood as a *sequel* resource and topic: for those who want to learn more in an industrial perspective. So while we’re about to list a number of deficiencies inherent in the code in fig. 1, it’s *not* to its detriment.

- The operation is blocking: that is, it will run until termination. An alternative is some functionality that allows for partial processing, after which control is returned to the caller (which may, say, update a progress bar) which can then resume the next “chunk” of processing. This plays a particular role in stability: if you give your program a blank check for as much memory needed to decompress something, you may end up using a *lot* of memory. So it’s nice to say “please decompress at most *X* bytes”, or thereabouts.
- More generally there is the notion of “library support”. We do have compression algorithms as stand-alone utilities, but you can also imagine a larger piece of software wanting to compress/decompress things in the course of its execution. The previous point is just one example: it may be nice to expose settings and behaviors programmatically, for other software packages to use.

¹The URL is here

```

1 private static final int R = 256;           // number of input chars
2 private static final int L = 4096;         // number of codewords = 2^W
3 private static final int W = 12;          // codeword width
4
5 // snip
6
7 public static void compress() {
8     String input = BinaryStdIn.readString();
9     TST<Integer> st = new TST<Integer>();
10
11     // since TST is not balanced, it would be better to insert in a different order
12     for (int i = 0; i < R; i++)
13         st.put("" + (char) i, i);
14
15     int code = R+1; // R is codeword for EOF
16
17     while (input.length() > 0) {
18         String s = st.longestPrefixOf(input); // Find max prefix match s.
19         BinaryStdOut.write(st.get(s), W);    // Print s's encoding.
20         int t = s.length();
21         if (t < input.length() && code < L) // Add s to symbol table.
22             st.put(input.substring(0, t + 1), code++);
23         input = input.substring(t);          // Scan past s in input.
24     }
25     BinaryStdOut.write(R, W);
26     BinaryStdOut.close();
27 }

```

Figure 1: A Textbook Implementation of LZW-Compress

- There is no debugging support. While the conceit is that this is a “finished” algorithm and presumably once we’re confident it its correctness we wouldn’t do ongoing development, in practice that rarely happens. We may want to add more logging, we may want to insert special hardware-specific subroutines, or something else. Having trace statements, assertions, and similar is important and keeps future development, tractable.

The LZW implementation we’ll develop here will have the following properties:

1. Variable-width encodings.
2. Clear-codes, in which the encoder can put a distinguished key in the stream that tells the dictionary to reset to the default.
3. An interface design—this influences the `encode` and `decode`—that takes into account some practical applications (memory usage, etc.).
4. A harness of the library into a simple CLI program that can be used to compress, decompress, etc.

```

1  uint64_t bitwrite_buffer = 0;
2  uint32_t bitwrite_buffer_size = 0;
3  const uint32_t BITWRITE_BUFFER_MAX_SIZE = sizeof(bitwrite_buffer) * 8;
4
5  void bitwrite_buffer_push_bits(uint32_t v, uint8_t l) {
6      ASSERT(bitwrite_buffer_size + l < BITWRITE_BUFFER_MAX_SIZE);
7      uint32_t mask = (1 << l) - 1;
8      bitwrite_buffer = (bitwrite_buffer << l) | (v & mask);
9      bitwrite_buffer_size += l;
10 }
11
12 uint8_t bitwrite_buffer_pop_byte(void) {
13     ASSERT(bitwrite_buffer_size >= 8);
14     bitwrite_buffer_size -= 8;
15     uint8_t b = bitwrite_buffer >> bitwrite_buffer_size;
16     return b;
17 }

```

Figure 2: Bit Write Buffer Implementation

5. Extensive fuzzing and performance testing, and consequent (slight) performance focus.

The core implementation is under 600 lines of C code, and corresponding driver less than 300, totalling to fewer than 1,000 lines of code for the full implementation.

3. Implementation Details

3.1. Reading and Writing Bits and Bytes

We can typically only write bytes (that is, 8 bits) to a file at a time. If we encode a stream into 103 keys, each key being 9 bits long, we’ve set ourselves up for emitting 927 bits—obviously not neatly fitting into some number of bytes.

Solving this problem can be a motivating introduction to bitwise tricks and using those more “exotic” operators like `&`, `|`, and `<<`. We shall see that, with these tools, we can implement a way of *buffering* up unusual-length-values into a 64-bit value, and then “peel off” 8-bit chunks from that buffer as needed. The complete implementation is in fig. 2.

Our actual “buffer”, `bitwrite_buffer`, is simply a 64-bit integer. We’ll see over the course of our implementation that this basically restricts our key-length to 32-bit integers.

Digression: Bounded-Size Buffers At the onset of my programming journey I would be very uncomfortable with this: having implicit resource limits felt very “vulnerable”. A program hitting a resource limit (such as, in this case, secretly trying to encode too many bits at once in our buffer) can lead to hard-to-investigate bugs and a lot of pain.

```

1  #ifndef IO_BUFFER_SIZE
2  #define IO_BUFFER_SIZE 4096
3  #endif
4  static uint8_t fwrite_buffer[IO_BUFFER_SIZE];
5  static int emit_buffer_next = 0;
6  void write_buffer_flush() {
7      fwrite(fwrite_buffer, 1, emit_buffer_next, lzw_output_file);
8      emit_buffer_next = 0;
9  }
10 void lzw_write_byte(uint8_t c) {
11     if (emit_buffer_next == sizeof(fwrite_buffer)) {
12         write_buffer_flush();
13     }
14     fwrite_buffer[emit_buffer_next++] = c;
15 }

```

Figure 3: Emitting Bytes

So I would have gone to great pains to make some kind of resizable bit-buffer that could handle any number of bits.

However, with experience I improved my intuition about which resource limits would actually be hit in development, and so I could feel more confident about “putting things off” for later. I also became more experienced with (or inured to?) investigating weird bugs and fixing them—still not something I sought out, but no longer a terror-in-the-dark. Lastly, you can see the most critical tool used here: `ASSERT` statements that will fail in the case the resource is exhausted. The flexibility to define “solving” the problem as “telling the developer about it, and failing” is a remarkably powerful one. Of course for commercial software it’s unsatisfying, but for internal-facing projects like this it’s great.

A maybe-unintuitive property is that we use the same emit bytebuffer in both the decode and encode cases. Maybe we can do something more optimally, then, in the decode/encode case, as we know that one direction of each isn’t going to be variable-width.

Buffered IO This pattern of “empty the buffer when we can’t fit the next thing” is pretty common. A common follow-on bug is forgetting to flush the last thing.

Why use a `#define`? Partly for nostalgia-esque old-fashioned, I suppose. A lot of runtime options or flags or configuration can instead be compile-time. You can imagine, if we were allocating `fwrite_buffer` on the heap, we’d pass in the size as a runtime parameter. Here, however, we want to reserve the space for the buffer in the process-level memory, so we can pass the size as a *build* option: `gcc lzw.c -DIO_BUFFER_SIZE 10000`, for instance.

Why have a buffer at all? Semantically, we don’t need the buffer at all: we can just call `fwrite` (or `fputc`) directly, one byte at a time. There is the general intuition that

doing things in batches is “nicer” for computers and leads to better performance: indeed on my machine it lead to a considerable improvement.

It’s worth considering *why* (as in, what is the actual mechanism that lead to a speedup) in such a case. My initial hypothesis was that each `fputc` is a system call, and crossing that boundary is expensive. However, using `strace` (a utility that logs every system call) shows that the system itself *does* have a buffer (which is not too surprising) and so we actually have the same number of system calls.

This is a common challenge I’ve seen in performance work: people take the success of an optimization as proof that they’re hypothesis is write: for this situation I don’t know what actual mechanism changed to lead to a performance improvement, only that it’s considerable! This is a great homework problem, maybe. If I wanted to spend more time on this I’d next hypothesis something about dynamic linking or indirect calls, but again I don’t know.

Any gotchas? Yes, in particular we have to remember to flush this buffer, see what was do in `lzw_encode_end` in the listing at the end.

3.2. Interlude: Modules versus Objects

The state of our program is all globals. This is weird—why do we we do this? It avoid heap allocations, which are nice, at the cost that you can only do this one at a time. It’s not that bad—moreover, we can do more things compile-time.

3.3. State-Machine Based Dictionary

In the common internet example the dictionary is more-or-less literal. The TST object Sedgwick uses is some special-purpose tree that takes the head of the input remainder and will search it for as long as possible. To me this doesn’t feel natural (though maybe his approach is the better one, ultimately!) because there’s an aspect to how the `longestPrefixOf` method is what’s (behind the call-site) consuming the input.

(In particular, this doesn’t lend to a natural API extension that permits the user to specify “only encode the first N bytes”.)

Figure 4 shows how this code reads in input for encoding (compression), and in fig. 5 we see how taking the next bit (`c`) updates our `curr` node to either the next already-existing child in our tree, or allocate the new one.

So the dictionary is pretty important. Intuitively, that’s the thing we’re really generating, and using, a lot.

3.4. Data Structure Design

The structure is pretty exotic. Behold this diagram:

Short String Optimization We save a lot of memory with that crazy union thing. This is definitely a case where we have to decide on a heuristic: in industry figuring out

```

1 size_t lzw_encode(size_t l) {
2     for (;;) {
3         int c = lzw_read_byte();
4         if (c == EOF) {
5             lzw_encode_end();
6             break;
7         }
8         if (lzw_next_char(c) != NEXT_CHAR_CONTINUE) // begin processing...

```

Figure 4: The control-flow driving compression

```

1 int lzw_next_char(uint8_t c) {
2     lzw_node_p next = children_set_find(&curr->children, c);
3     if (next) {
4         curr = next;
5         return NEXT_CHAR_CONTINUE;
6     }
7     if (lzw_max_key && lzw_next_key >= lzw_max_key) {
8         DTRACE(DB_STATE, "APPEND(0x%x)\tMAXKEY\n", c);
9         return NEXT_CHAR_MAX;
10    }
11    // Else, allocate a new leaf in our tree.
12    // Note curr is NOT updated: we still need to emit the "old" prefix in this case.
13    const uint32_t k = lzw_next_key++;
14    next = children_set_allocate(&curr->children, c, k);
15    lzw_data[k].parent = next;
16    return NEXT_CHAR_NEW;
17 }

```

Figure 5: Evolving our dictionary object

the right way to “define success” (including how to update that definition!) is the most important thing.

3.5. Feature: Clear Codes

This was a bear to get the coordination right: I hadn’t anticipated the subtle issue about increasing the length as if it’s a new key.

It’s also weird that it counts as bytes we encode/decode.

3.6. Using our Features: EMA

Learned that trick from my SAT solver. It’s actually really nice! From my motivating input it definitely worked. Dream: avoid having to emit clear codes at all. I was really close, but coordinating the implicit state was even worse. Maybe that’s a homework problem?

3.7. Debugging Support

See the trace statements? Articulating the streams as emitting keys, and emitting logical bits, and then emitted physical bytes, was very useful for debugging. Supplementary tools: parse and compose log statements.

4. Testing: Fuzzing

Fuzzing was a lot of fun! It was more a quick way to determine that nothing regressed, but also found some awful edge-cases with clear-codes. Frustration: lack of tooling for futzing with command line parameters, basically.

5. Performance

I don’t get hardware counters. There’s encoding speed-and-memory-usage, decoding speed-and-memory-usage, and compression ratio. With some futzing I did “ok” compared to gzip on our motivating input. Still much slower, not sure why: presumably data structure is just bad, but would want hardware counters.

Attribution Problem in Performance I worked on a project where our software supported multiple workloads, call them A, B, C, etc. One workload, P, was both *extremely popular* and also *very resource intensive*. A single process could be working on multiple workloads at once. A common refrain I had to push back against was “oh, every single crash/resource-exhaustion is attributable to P”. On the one hand that’s a reasonable starting point, but tautologically *every* trace of a process is going to include a P: it was by far the most common one. I’m sure there’s some cutesy statistical phrasing for this behavior, but this was a surprising example of heuristics biting people—often the

issue was some other, less-exercised workload. (Nevermind the higher-level issues of load-balancing etc.)

5.1. Compression Ratio

This is trying to be the “textbook definition”, so except up to EMA stuff I didn’t expect much improvement. That EMA really helps is very cool.

5.2. Throughput

Need hardware counters! I can continue to guess what’s wrong, but it’s just not productive (or satisfying).

6. What’s Next?

- Getting hardware counters to help inform performance work.
- It’s not clear if, e.g., setting the max key to 1024 would basically induce an off-by-one in our max-length (we should only need 10 bits for each symbol, but we may accidentally get up to 11).
- I’d like to actually use this as a library for another project. That may mean making it more “ergonomic” (and better) to use the EMA stuff (as in, provide it for those users). It may also, sadly, mean that I should change this from a module to an object, so we actually have multiple such threads going on at once, haha.

7. Industrial Comparison

A homework assignment is to compare this to a real compression algorithm.

A. Core Library Implementation

```
1  #include <assert.h>
2  #include <stdbool.h>
3  #include <stdint.h>
4  #include <stdio.h>
5  #include <stdlib.h>
6  #include <string.h>
7  #include <unistd.h>
8
9  #include "lzw.h"
10
11 typedef struct lzw_node_tag *lzw_node_p;
12
13 typedef struct lzw_children_set_tag {
```

```

14     bool use_array;
15     struct {
16         int index;
17         struct {
18             uint8_t key;
19             lzw_node_p value;
20         } immediate[4];
21     } local;
22     lzw_node_p *all;
23 } lzw_children_set_t;
24
25 // we want to build a mapping from keys to data-strings.
26 typedef struct lzw_node_tag {
27     lzw_children_set_t children;
28     uint32_t key;
29 } lzw_node_t, *lzw_node_p;
30
31 typedef struct {
32     uint8_t *data;
33     uint32_t len;
34     lzw_node_p parent;
35 } lzw_data_t;
36
37 lzw_data_t *lzw_data = NULL;
38 lzw_node_p root = NULL;
39 lzw_node_p curr = NULL;
40
41 uint32_t lzw_length = 0;
42 uint32_t lzw_next_key = 0;
43 uint32_t lzw_max_key = 0;
44
45 FILE *lzw_input_file;
46 FILE *lzw_output_file;
47
48 uint64_t lzw_bytes_written = 0;
49 uint64_t lzw_bytes_read = 0;
50
51 // There's an implicit invariant that our lzw_length can never be more than 32,
52 // and this means our buffers always need at least 2 uses before overflowing.
53 // Some of our iteration depends on that (in particular write_key, which
54 // unconditionally enbuffers something before trying to drain it).
55 uint64_t bitread_buffer = 0;
56 uint32_t bitread_buffer_size = 0;
57 const uint32_t BITREAD_BUFFER_MAX_SIZE = sizeof(bitread_buffer) * 8;
58 uint64_t bitwrite_buffer = 0;
59 uint32_t bitwrite_buffer_size = 0;
60 const uint32_t BITWRITE_BUFFER_MAX_SIZE = sizeof(bitwrite_buffer) * 8;
61
62 // Before we get too much into executable code,

```

```

63 // we want to express the different modes we can run in.
64 // I.e., debug levels
65 #ifdef NDEBUG
66 #define DTRACE(k, ...)
67 #define ASSERT(x)
68 #define DEBUG_STMT(x)
69 #else
70 enum {
71     DB_STATE,
72     DB_BYTE_STREAM,
73     DB_KEY_STREAM,
74     DB_DICTIONARY,
75     DB_MAX,
76 };
77 int DB_KEYS_SET[DB_MAX] = {};
78 #define DTRACE(k, ...)
79     if (DB_KEYS_SET[k]) {
80         fprintf(stderr, __VA_ARGS__);
81     }
82 #define ASSERT(x) assert(x)
83 #define DEBUG_STMT(x) x
84 #endif
85
86 void lzw_set_debug_string(const char *s) {
87     #ifdef NDEBUG
88     #else
89         for (int i = 0; i < strlen(s); i++) {
90             switch (s[i]) {
91                 case 's':
92                     DB_KEYS_SET[DB_STATE] = 1;
93                     break;
94                 case 'b':
95                     DB_KEYS_SET[DB_BYTE_STREAM] = 1;
96                     break;
97                 case 'k':
98                     DB_KEYS_SET[DB_KEY_STREAM] = 1;
99                     break;
100                 case 'd':
101                     DB_KEYS_SET[DB_DICTIONARY] = 1;
102                     break;
103             }
104         }
105     #endif
106 }
107
108 #ifndef NDEBUG
109 static const char *asbits(uint64_t v, uint8_t l) {
110     static char format[(1 << 8 * sizeof(uint8_t)) + 1];
111     // Emit bits backwards

```

```

\
\
\

```

```

112     for (int i = 0; i < l; i++) {
113         format[l - (i + 1)] = v % 2 ? '1' : '0';
114         v >>= 1;
115     }
116     format[l] = '\0';
117     return format;
118 }
119 #endif
120
121 lzw_node_p children_set_find(lzw_children_set_t *s, uint8_t k) {
122     if (!s->use_array) {
123         int max_index = s->local.index;
124         for (int i = 0; i < max_index; i++) {
125             if (s->local.immediate[i].key == k) {
126                 return s->local.immediate[i].value;
127             }
128         }
129         return NULL;
130     }
131     return s->all[k];
132 }
133
134 lzw_node_p children_set_allocate(lzw_children_set_t *s, uint8_t c, uint32_t k) {
135     lzw_node_p r = calloc(1, sizeof(lzw_node_t));
136     r->key = k;
137     if (!s->use_array) {
138         int n = s->local.index++;
139         if (n < 4) {
140             s->local.immediate[n].key = c;
141             s->local.immediate[n].value = r;
142             return r;
143         }
144         s->use_array = true;
145         s->all = calloc(256, sizeof(lzw_node_p));
146         for (int i = 0; i < n; i++) {
147             s->all[s->local.immediate[i].key] = s->local.immediate[i].value;
148         }
149     }
150     s->all[c] = r;
151     return r;
152 }
153
154 enum { NEXT_CHAR_CONTINUE = 0, NEXT_CHAR_MAX = 1, NEXT_CHAR_NEW = 2 };
155 const uint32_t lzw_clear_code = 256;
156
157 // The primary action of this table is to ingest
158 // the next byte, and maintain the correct encoding
159 // information for the implicit string seen-so-far.
160 // That's captured in this function:

```

```

161 int lzw_next_char(uint8_t c) {
162     lzw_node_p next = children_set_find(&curr->children, c);
163     if (next) {
164         DTRACE(DB_STATE, "APPEND(%#x)\t\tSTATE\t%u\t%u\n", c, curr->key, next->key);
165         curr = next;
166         return NEXT_CHAR_CONTINUE;
167     }
168     // we have reached the end of the string.
169     if (lzw_max_key && lzw_next_key >= lzw_max_key) {
170         DTRACE(DB_STATE, "APPEND(%#x)\tMAXKEY\n", c);
171         return NEXT_CHAR_MAX;
172     }
173     // Create the new fields for the new node
174     const uint32_t k = lzw_next_key++;
175     next = children_set_allocate(&curr->children, c, k);
176     const uint32_t l = curr->key == -1 ? 0 : lzw_data[curr->key].len;
177     uint8_t *data = calloc(l + 1, sizeof(uint8_t));
178     if (l) {
179         memcpy(data, lzw_data[curr->key].data, l * sizeof(uint8_t));
180     }
181     data[l] = c;
182     lzw_data[k].data = data;
183     lzw_data[k].len = l + 1;
184     lzw_data[k].parent = next;
185     DTRACE(DB_DICTIONARY, "DICT\tADD\t%u\t%u\t%u\t%#x\n", k, curr->key, l + 1, c);
186     DTRACE(DB_STATE, "APPEND(%#x)\t\tNEWSTATE %u\n", c, k);
187     return NEXT_CHAR_NEW;
188 }
189
190 // We also have book-keeping of when we have to
191 // update the length
192 static size_t lzw_data_size(void) {
193     return (1 << lzw_length) * sizeof(lzw_data_t);
194 }
195
196 void lzw_len_update() {
197     DTRACE(DB_STATE, "INLENGTH %d->%d\n", lzw_length, lzw_length + 1)
198     size_t old_length = lzw_data_size();
199     lzw_length++;
200     size_t new_length = lzw_data_size();
201     ASSERT(old_length * 2 == new_length);
202     lzw_data = realloc(lzw_data, new_length);
203     memset((char *)lzw_data + old_length, 0, old_length);
204 }
205
206 void free_dictionary(lzw_node_p t) {
207     if (!t)
208         return;
209     if (t->children.use_array) {

```

```

210     for (uint16_t i = 0; i < 256; ++i) {
211         free_dictionary(t->children.all[i]);
212     }
213     free(t->children.all);
214 } else {
215     int n = t->children.local.index;
216     for (int i = 0; i < n; i++) {
217         free_dictionary(t->children.local.immediate[i].value);
218     }
219 }
220 free(t);
221 }
222
223 void lzw_destroy_state(void) {
224     free_dictionary(root);
225     curr = NULL;
226     root = NULL;
227     for (uint32_t i = 0; i < lzw_next_key; ++i) {
228         if (lzw_data[i].data)
229             free(lzw_data[i].data);
230     }
231     if (lzw_data) {
232         free(lzw_data);
233         lzw_data = NULL;
234     }
235     lzw_next_key = 0;
236     ASSERT((bitread_buffer & ((1 << bitread_buffer_size) - 1)) == 0);
237     ASSERT(bitwrite_buffer_size == 0);
238 }
239
240 void bitwrite_buffer_push_bits(uint32_t v, uint8_t l) {
241     ASSERT(bitwrite_buffer_size + 1 < BITWRITE_BUFFER_MAX_SIZE);
242     uint32_t mask = (1 << l) - 1;
243     bitwrite_buffer = (bitwrite_buffer << l) | (v & mask);
244     bitwrite_buffer_size += l;
245 }
246
247 uint8_t bitwrite_buffer_pop_byte(void) {
248     ASSERT(bitwrite_buffer_size >= 8);
249     bitwrite_buffer_size -= 8;
250     uint8_t b = bitwrite_buffer >> bitwrite_buffer_size;
251     return b;
252 }
253
254 #ifndef IO_BUFFER_SIZE
255 #define IO_BUFFER_SIZE 4096
256 #endif
257 static uint8_t fwrite_buffer[IO_BUFFER_SIZE];
258 static int emit_buffer_next = 0;

```

```

259 void write_buffer_flush() {
260     fwrite(fwrite_buffer, 1, emit_buffer_next, lzw_output_file);
261     emit_buffer_next = 0;
262 }
263 void lzw_write_byte(uint8_t c) {
264     if (emit_buffer_next == sizeof(fwrite_buffer)) {
265         write_buffer_flush();
266     }
267     fwrite_buffer[emit_buffer_next++] = c;
268 }
269
270 // Reading v from "left to right", we
271 // emit the l bits of v.
272 void write_key(uint32_t v, uint8_t l) {
273     ASSERT((v & ((1 << l) - 1)) == v); // v doesn't have extra bits
274     DTRACE(DB_KEY_STREAM, "EMITKEY(%d):\t\t%d\t\t%s\n", l, v, asbits(v, l));
275     bitwrite_buffer_push_bits(v, l);
276     while (bitwrite_buffer_size >= 8) {
277         uint8_t c = bitwrite_buffer_pop_byte();
278         DTRACE(DB_BYTE_STREAM, "EMITBYTE(encode):\t\t%x\t\t%s\n", c, asbits(c, 8));
279         lzw_write_byte(c);
280         lzw_bytes_written++;
281     }
282 }
283
284 bool key_requires_bigger_length(uint32_t k) { return k >= (1 << lzw_length); }
285 bool update_length() {
286     if (key_requires_bigger_length(lzw_next_key)) {
287         lzw_len_update();
288         return true;
289     }
290     return false;
291 }
292
293 void lzw_init() {
294     root = (lzw_node_p) calloc(1, sizeof(lzw_node_t));
295     curr = root;
296     root->key = -1;
297     lzw_length = 1;
298     lzw_next_key = 0;
299     lzw_data = calloc(1 << lzw_length, sizeof(lzw_data_t));
300
301     #ifndef NDEBUG
302     int old_state_key = DB_KEYS_SET[DB_STATE];
303     DB_KEYS_SET[DB_STATE] = 0;
304     int old_dict_key = DB_KEYS_SET[DB_DICTIONARY];
305     DB_KEYS_SET[DB_DICTIONARY] = 0;
306     #endif
307     for (uint16_t i = 0; i < 256; ++i) {

```



```

308     lzw_next_char(i);
309     update_length();
310 }
311 #ifndef NDEBUG
312     DB_KEYS_SET[DB_STATE] = old_state_key;
313     DB_KEYS_SET[DB_DICTIONARY] = old_dict_key;
314 #endif
315
316     ASSERT(lzw_clear_code == lzw_next_key);
317     lzw_next_key++; // reserve 256 for the clear-code.
318     update_length();
319
320     bitread_buffer = 0;
321     bitread_buffer_size = 0;
322
323     bitwrite_buffer = 0;
324     bitwrite_buffer_size = 0;
325
326     lzw_bytes_read = 0;
327     lzw_bytes_written = 0;
328 }
329
330 static uint8_t fread_buffer[IO_BUFFER_SIZE];
331 static int read_buffer_next = 0;
332 static int read_buffer_max = 0;
333 uint32_t lzw_read_byte(void) {
334     if (read_buffer_next == read_buffer_max) {
335         read_buffer_max =
336             fread(fread_buffer, 1, sizeof(fread_buffer), lzw_input_file);
337         read_buffer_next = 0;
338     }
339     if (read_buffer_max == 0) {
340         return EOF;
341     }
342     return fread_buffer[read_buffer_next++];
343 }
344
345 bool input_eof(void) { return read_buffer_max == 0 && feof(lzw_input_file); }
346
347 size_t lzw_encode(size_t l) {
348     size_t i = 0;
349     for (;;) {
350         int c = lzw_read_byte();
351         if (c == EOF) {
352             DTRACE(DB_BYTE_STREAM, "READBYTE(encode):\t%d\n", c);
353             DTRACE(DB_STATE, "lzw_encode:eof\n");
354             lzw_encode_end();
355             break;
356         }

```

```

357     i++;
358     DTRACE(DB_BYTE_STREAM, "READBYTE(encode):\t%x\t%s\n", c, asbits(c, 8));
359     if (lzw_next_char(c) != NEXT_CHAR_CONTINUE) {
360         write_key(curr->key, lzw_length);
361         curr = root;
362         update_length();
363         lzw_next_char(c);
364         if (i > 1) {
365             break;
366         }
367     }
368 }
369 lzw_bytes_read += i;
370 return i;
371 }
372
373 void lzw_write_clear_code(void) {
374     DTRACE(DB_STATE, "CLEAR_CODE\t%zu\t%d\n", lzw_bytes_written, input_eof());
375     if (input_eof()) {
376         return; // don't bother
377     }
378     if (curr != root) {
379         write_key(curr->key, lzw_length);
380         curr = root;
381         // When we read in a code, we always assume that it's a new key
382         // (unless if we're at the max). So our reader preemptively
383         // updates the length at the key boundaries---we need to do that
384         // here too, even though this key isn't new.
385         if (key_requires_bigger_length(lzw_next_key + 1)) {
386             lzw_len_update();
387         }
388     }
389     write_key(lzw_clear_code, lzw_length);
390     lzw_encode_end();
391 }
392
393 void lzw_encode_end(void) {
394     // if we haven't done anything yet, make that more explicit
395     DTRACE(DB_STATE, "ENCODE_END\t%u\t%zu\t%d\n", bitwrite_buffer_size,
396         lzw_bytes_written, curr == root);
397     if (bitwrite_buffer_size == 0 && lzw_bytes_written == 0 && curr == root) {
398         return;
399     }
400     if (curr != root) {
401         write_key(curr->key, lzw_length);
402         curr = root;
403     }
404     if (bitwrite_buffer_size != 0) {
405         // We want to finish emitting our last key.

```

```

406     // cap off our buffer: there are (say) 3 valid bits left,
407     // we just need to pad it so we can emit those 3 bits as part
408     // of a larger byte.
409     uint8_t bits_to_add = 8 - (bitwrite_buffer_size % 8);
410     write_key(0, bits_to_add);
411     ASSERT(bitwrite_buffer_size == 0);
412 }
413 write_buffer_flush();
414 }
415
416 void bitread_buffer_push_byte(uint8_t c) {
417     ASSERT(bitread_buffer_size + 8 < BITREAD_BUFFER_MAX_SIZE);
418     bitread_buffer <<= 8;
419     bitread_buffer |= c;
420     bitread_buffer_size += 8;
421 }
422
423 uint32_t bitread_buffer_pop_bits(uint32_t bitcount) {
424     ASSERT(bitread_buffer_size >= bitcount);
425     uint64_t bitread_buffer_copy = bitread_buffer;
426     // slide down the "oldest" bits
427     bitread_buffer_copy >>= (bitread_buffer_size - bitcount);
428     bitread_buffer_copy &= (1 << bitcount) - 1;
429     bitread_buffer_size -= bitcount;
430     return bitread_buffer_copy;
431 }
432
433 // This will read the next bits up to our buffer.
434 bool read_bits(uint32_t *v) {
435     while (bitread_buffer_size < lzw_length) {
436         uint32_t c = lzw_read_byte();
437         if (c == EOF) {
438             DTRACE(DB_BYTE_STREAM, "READBYTE(decode):\t%d\n", c);
439             return false;
440         }
441         DTRACE(DB_BYTE_STREAM, "READBYTE(decode):\t%x\t%s\n", c, asbits(c, 8));
442         lzw_bytes_read++;
443         bitread_buffer_push_byte(c);
444     }
445     *v = bitread_buffer_pop_bits(lzw_length);
446     return true;
447 }
448
449 bool lzw_valid_key(uint32_t k) {
450     ASSERT(k < (1 << (lzw_length)));
451     return lzw_data[k].data != NULL;
452 }
453
454 size_t lzw_decode(size_t limit) {

```

```

455     uint32_t curr_key;
456     size_t read = 0;
457     while (read < limit && read_bits(&curr_key)) {
458         DTRACE(DB_KEY_STREAM, "READKEY(%d):\t\t%d\t%s\n", lzw_length, curr_key,
459             asbits(curr_key, lzw_length));
460         if (curr_key == lzw_clear_code) {
461             DTRACE(DB_STATE, "DECODE\tCLEAR_CODE\n");
462             lzw_destroy_state();
463             // Curious thing: we can return a value greater than lzw_bytes_read,
464             // as lzw_init() set that back to 0. We continue because we also
465             // promise to always emit something when we're called.
466             lzw_init();
467             continue;
468         }
469         ASSERT(lzw_valid_key(curr_key));
470
471         // emit that string:
472         uint8_t *s = lzw_data[curr_key].data;
473         uint32_t l = lzw_data[curr_key].len;
474         ASSERT(l);
475         for (uint32_t i = 0; i < l; ++i) {
476             DTRACE(DB_BYTE_STREAM, "EMITBYTE(decode):\t%x\n", s[i]);
477             lzw_write_byte(s[i]);
478         }
479         lzw_bytes_written += l;
480         read += l;
481         curr = lzw_data[curr_key].parent;
482
483         if (key_requires_bigger_length(lzw_next_key + 1)) {
484             lzw_len_update();
485         }
486
487         // peek at the next string:
488         DTRACE(DB_STATE, "DECODE(peek)\n");
489         if (!read_bits(&curr_key)) {
490             DTRACE(DB_STATE, "DECODE(break)\n");
491             // We're at EOF, so just early-out
492             break;
493         }
494         if (curr_key == lzw_clear_code) {
495             // Skip our checks: we don't want to evolve our state.
496             bitread_buffer_size += lzw_length;
497             continue;
498         }
499         DTRACE(DB_STATE, "DECODE(continue)\n");
500         bitread_buffer_size += lzw_length;
501
502         // Find the next character.
503         // If the next key is valid, that means

```

```

504     // the next key is already something we've seen,
505     // otherwise it's going to be the key for the new
506     // string. Either way, we take the next step
507     // on that character, but then manually reset
508     // our curr node back to the root in prep for
509     // really reading the next string.
510     if (lzw_valid_key(curr_key)) {
511         s = lzw_data[curr_key].data;
512     } else {
513         ASSERT(curr_key - 1 == lzw_clear_code || lzw_valid_key(curr_key - 1));
514     }
515     DEBUG_STMT(int b =)
516     lzw_next_char(s[0]);
517     ASSERT(b != NEXT_CHAR_CONTINUE);
518     curr = root;
519 }
520 write_buffer_flush();
521 return read;
522 }

```

B. Command Line Implementation

```

1  #include "lzw.h"
2  #include <assert.h>
3  #include <stdbool.h>
4  #include <stdint.h>
5  #include <stdio.h>
6  #include <stdlib.h>
7  #include <string.h>
8  #include <unistd.h>
9
10 int getopt(int, char *const[], const char *);
11 char *optarg;
12
13 bool do_decode = false;
14 bool do_encode = false;
15 bool do_ratio = false;
16 bool trace_ratio = false;
17 char *ratio_log_filename = NULL;
18
19 FILE *user_input;
20 FILE *user_output;
21
22 int verbosity = 0;
23
24 size_t page_size = 4096;
25

```

```

26  uint64_t total_stream_read = 0;
27  uint64_t total_stream_written = 0;
28
29  uint64_t prev_bytes_written = 0;
30  uint64_t prev_bytes_read = 0;
31  uint64_t page_bytes_written = 0;
32  uint64_t page_bytes_read = 0;
33  double next_ratio() {
34      page_bytes_written = lzw_bytes_written - prev_bytes_written;
35      page_bytes_read = lzw_bytes_read - prev_bytes_read;
36      prev_bytes_written = lzw_bytes_written;
37      prev_bytes_read = lzw_bytes_read;
38      // fprintf(stderr, "page_bytes_written: %zu\n", page_bytes_written);
39      // fprintf(stderr, "page_bytes_read: %zu\n", page_bytes_read);
40      assert(page_bytes_written);
41      double compression_ratio =
42          (double)page_bytes_written / (double)page_bytes_read;
43      if (do_decode) {
44          compression_ratio = (double)page_bytes_read / (double)page_bytes_written;
45      }
46      return compression_ratio;
47  }
48
49  void reset_written() {
50      prev_bytes_written = 0;
51      prev_bytes_read = 0;
52  }
53
54  void decode_stream() {
55      total_stream_written = 0;
56      lzw_init();
57      // Decode is guaranteed to make progress (even in presence of clear-codes)
58      for (;;) {
59          size_t written = lzw_decode(page_size);
60          if (!written) {
61              break;
62          }
63          total_stream_written += written;
64      }
65      lzw_destroy_state();
66  }
67
68  void encode_stream() {
69      total_stream_read = 0;
70      total_stream_written = 0;
71      FILE *ratio_log_file = stderr;
72      if (ratio_log_filename) {
73          ratio_log_file = fopen(ratio_log_filename, "w");
74      }

```

```

75
76  const int ema_delay = 32;
77
78  for (int block_count = 0;; block_count++) {
79      reset_written();
80      double ema_slow = 0.0;
81      double ema_slow_alpha = 0.0001;
82      double ema_fast = 0.0;
83      double ema_fast_alpha = 0.01;
84
85      lzw_init();
86
87      for (int page_count = 0;; page_count++) {
88          // fprintf(stderr, "processing page: %d\n", page_count);
89          size_t bytes_processed = lzw_encode(page_size);
90          // We've hit EOF.
91          if (!bytes_processed) {
92              total_stream_read += lzw_bytes_read;
93              total_stream_written += lzw_bytes_written;
94              lzw_destroy_state();
95              return;
96          }
97
98          // We've processed a page's worth of data, now
99          // evaluate our compression ratio and windows.
100         double compression_ratio = next_ratio();
101         if (page_count < ema_delay) {
102             ema_slow = compression_ratio;
103             ema_fast = compression_ratio;
104         } else {
105             ema_slow += ema_slow_alpha * (compression_ratio - ema_slow);
106             ema_fast += ema_fast_alpha * (compression_ratio - ema_fast);
107         }
108         if (trace_ratio) {
109             fprintf(ratio_log_file,
110                 "%s ratio: %f\tema_slow=%f\tema_fast=%f\tpage_bytes_read: "
111                 "%zu\tpage_bytes_written: %zu\n",
112                 do_encode ? "compression" : "decompression",
113                 compression_ratio, ema_slow, ema_fast, page_bytes_read,
114                 page_bytes_written);
115         }
116
117         // Now consume our ratio information: should we start a new block?
118         if (do_ratio && page_count >= ema_delay &&
119             (ema_slow * 1.5 < ema_fast || ema_fast > 0.7)) {
120             if (trace_ratio) {
121                 fprintf(ratio_log_file, "resetting %d\n", page_count);
122             }
123             lzw_write_clear_code();

```

```

124         break; // this will lead to the destroy-state and init on the back-edge
125     }
126 }
127 total_stream_read += lzw_bytes_read;
128 total_stream_written += lzw_bytes_written;
129 lzw_destroy_state();
130 }
131 }
132
133 // process_stream consumes all the globally-set parameters
134 void process_stream() {
135     if (do_decode) {
136         decode_stream();
137     } else {
138         encode_stream();
139     }
140 }
141
142 // This is a shared correctness routine/helper
143 void init_streams(char *inbuffer, size_t insize, char **outbuffer,
144                  size_t *outsize) {
145     FILE *out = open_memstream(outbuffer, outsize);
146     FILE *in = fmemopen(inbuffer, insize, "r");
147     lzw_input_file = in;
148     lzw_output_file = out;
149 }
150 void close_streams(void) {
151     fclose(lzw_input_file);
152     fclose(lzw_output_file);
153 }
154
155 void round_trip() {
156     assert(user_output == stdout);
157     user_output = fopen("roundtrip_result.dat", "w");
158     assert(user_output);
159     do_encode = true;
160     do_decode = false;
161     FILE *intermediate = tmpfile();
162     lzw_output_file = intermediate;
163     fprintf(stderr, "Encoding stream\n");
164     process_stream();
165     fflush(intermediate);
166     fseek(intermediate, 0, SEEK_SET);
167     lzw_input_file = intermediate;
168     lzw_output_file = user_output;
169
170     do_encode = false;
171     do_decode = true;
172     fprintf(stderr, "Decoding stream\n");

```



```

173     process_stream();
174
175     fclose(intermediate);
176 }
177
178 void dumpbytes(const char *d, size_t c) {
179     for (size_t i = 0; i < c; i++) {
180         if (i > 0 && i % 20 == 0) {
181             fprintf(stderr, "\n");
182         }
183         fprintf(stderr, "%02x ", (unsigned char)d[i]);
184     }
185     fprintf(stderr, "\n");
186 }
187
188 void round_trip_in_memory(const char *Data, size_t Size) {
189     char *encodechunks = NULL;
190     size_t encodechunks_size = 0;
191     init_streams((char *)Data, Size, &encodechunks, &encodechunks_size);
192     do_encode = true;
193     do_decode = false;
194     process_stream();
195     close_streams();
196     assert(total_stream_read == Size);
197     assert(total_stream_written == encodechunks_size);
198
199     char *decodechunks = NULL;
200     size_t decodechunks_size;
201     init_streams(encodechunks, encodechunks_size, &decodechunks,
202                 &decodechunks_size);
203     do_encode = false;
204     do_decode = true;
205     process_stream();
206     close_streams();
207     // assert(total_stream_read == encodechunks_size); // because of clear-codes, this isn't quite
208     assert(total_stream_written == decodechunks_size); // We can compute this by summing decode ret
209
210     assert(Size == decodechunks_size);
211     assert(!memcmp(decodechunks, Data, Size));
212
213     free(encodechunks);
214     free(decodechunks);
215 }
216
217 #ifdef FUZZ_MODE
218 int LLVMFuzzerTestOneInput(const uint8_t *Data, size_t Size) {
219     if (Size < 12) { // just reserve enough bytes.
220         return 0;
221     }

```

```

222 // Peel off of the first chunk of data to choose our settings
223 lzw_max_key = *(uint32_t*)&Data[0];
224 if (lzw_max_key < 257) {
225     return 0;
226 }
227 page_size = *(uint32_t*)&Data[4];
228 if (page_size == 0) {
229     return 0;
230 }
231 do_ratio = (Data[9] % 2) == 0;
232 //fprintf(stderr, "page_size=%zu\tlzw_max_key=%u\tdo_ratio=%d\n", page_size, lzw_max_key, do_ratio);
233 round_trip_in_memory((const char *)Data+10, Size-10);
234 return 0;
235 }
236 #else
237 int main(int argc, char *argv[]) {
238     char c;
239     bool correctness_roundtrip = false;
240     bool correctness_roundtrip_memory = false;
241
242     user_input = stdin;
243     user_output = stdout;
244
245     while ((c = getopt(argc, argv, "deg:m:p:r:q:l:v:xCb:i:o:")) != -1) {
246         switch (c) {
247             case 'd':
248                 do_decode = true;
249                 break;
250             case 'e':
251                 do_encode = true;
252                 break;
253             case 'g':
254                 lzw_set_debug_string(optarg);
255                 break;
256             case 'm':
257                 lzw_max_key = atoi(optarg);
258                 break;
259             case 'p':
260                 page_size = atoi(optarg);
261                 break;
262             case 'q':
263                 trace_ratio = true;
264                 // if optarg is anything but '-'
265                 if (strcmp(optarg, "-")) {
266                     ratio_log_filename = strdup(optarg);
267                 }
268                 break;
269             case 'v':
270                 verbosity = atoi(optarg);

```

```

271         break;
272     case 'x':
273         do_ratio = true;
274         break;
275     case 'c': // for "correctness"
276         correctness_roundtrip = true;
277         break;
278     case 'C': // for "correctness"
279         correctness_roundtrip_memory = true;
280         break;
281     case 'i':
282         user_input = fopen(optarg, "r");
283         assert(user_input);
284         break;
285     case 'o':
286         user_output = fopen(optarg, "wx");
287         assert(user_output);
288         break;
289     default:
290         break;
291 }
292 }
293
294 if (correctness_roundtrip && correctness_roundtrip_memory) {
295     printf("Error, can't do both in-memory and through-file roundtrip (cC)\n");
296     return 2;
297 }
298 if (lzw_max_key && lzw_max_key < 256) {
299     printf("Error, max key too small (need >= 256, got %u)\n", lzw_max_key);
300     return 2;
301 }
302
303 if (trace_ratio && !do_ratio) {
304     fprintf(stderr,
305             "Warning, do_ratio=%s, trace_ratio=%s, unexpected behavior\n",
306             do_ratio ? "true" : "false", trace_ratio ? "true" : "false");
307 }
308
309 if (verbosity) {
310     fprintf(stderr, "lzw_max_key: %d\n", lzw_max_key);
311     fprintf(stderr, "page_size : %zu\n", page_size);
312 }
313
314 lzw_input_file = user_input;
315 lzw_output_file = user_output;
316
317 if (correctness_roundtrip) {
318     round_trip();
319 } else if (correctness_roundtrip_memory) {

```

```

320     char *inputbuffer = NULL;
321     size_t inputbuffer_size = 0;
322     FILE *copy_file = open_memstream(&inputbuffer, &inputbuffer_size);
323     char buffer[1028 * 1028];
324     size_t n = 0;
325     while ((n = fread(buffer, 1, sizeof(buffer), lzw_input_file)) > 0) {
326         fwrite(buffer, 1, n, copy_file);
327     }
328     fclose(lzw_input_file);
329     fclose(copy_file);
330     round_trip_in_memory(inputbuffer, inputbuffer_size);
331     free(inputbuffer);
332 } else {
333     if (do_encode == do_decode) {
334         printf("Error, must uniquely choose encode or decode\n");
335         return 1;
336     }
337     process_stream();
338 }
339
340 return 0;
341 }
342 #endif

```
